



Deep Learning School

Физтех-Школа Прикладной математики и информатики (ФПМИ) МФТИ

▼ Задание 3

Классификация текстов

В этом задании вам предстоит попробовать несколько методов, используемых в задаче классификации, а также понять насколько хорошо модель понимает смысл слов и какие слова в примере влияют на результат.

```

1 import pandas as pd
2 import numpy as np
3 import torch
4
5 from torchtext.legacy import datasets
6
7 from torchtext.legacy.data import Field, LabelField
8 from torchtext.legacy.data import BucketIterator
9
10 from torchtext.vocab import Vectors, GloVe
11
12 import torch.nn as nn
13 import torch.nn.functional as F
14 import torch.optim as optim
15 import random
16 from tqdm.autonotebook import tqdm

```

В этом задании мы будем использовать библиотеку torchtext. Она довольно проста в использовании и поможет нам сконцентрироваться на задаче, а не на написании DataLoader-а.

```

1 TEXT = Field(sequential=True, lower=True, include_lengths=True) # Поле текста
2 LABEL = LabelField(dtype=torch.float) # Поле метки

```

```

1 SEED = 1234
2
3 np.random.seed(SEED)
4 torch.manual_seed(SEED)
5 torch.random.manual_seed(SEED)
6 torch.cuda.random.manual_seed_all(SEED)
7 torch.backends.cudnn.deterministic = True

1 def calc_f1_score(model, iter, output_calc_func):
2     model.eval()
3
4     tp = 0.0
5     tn = 0.0
6     fp = 0.0
7     fn = 0.0
8     for batch in iter:
9         with torch.no_grad():
10             outputs = output_calc_func(model, batch)
11             outputs = (torch.sigmoid(outputs).cpu() > 0.5).int().squeeze(1)
12             labels = batch.label.cpu().int()
13
14             tp += (labels * outputs).sum()
15             tn += ((1 - labels) * (1 - outputs)).sum()
16             fp += ((1 - labels) * outputs).sum()
17             fn += (labels * (1 - outputs)).sum()
18
19     epsilon = 1e-7
20     precision = tp / (tp + fp + epsilon)
21     recall = tp / (tp + fn + epsilon)
22
23     f1_score = 2 * (precision * recall) / (precision + recall + epsilon)
24     f1_score = f1_score.item()
25
26     return f1_score

```

Напишем функции для обучения моделей

```

1 def train_rnn(model, loss_func, train_iter, val_iter, max_epochs, patience, max_grad_norm=2):
2     min_loss = np.inf
3
4     cur_patience = 0
5
6     for epoch in range(1, max_epochs + 1):
7         train_loss = 0.0
8         model.train()
9         pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
10        pbar.set_description(f"Epoch {epoch}")
11        for it, batch in pbar:
12            opt.zero_grad()
13
14            outputs = model(batch.text[0], batch.text[1].cpu())
15            loss = loss_func(outputs, batch.label.unsqueeze(1))
16            loss.backward()
17            if max_grad_norm is not None:
18                torch.nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
19            opt.step()
20
21            train_loss += loss.cpu().detach()
22
23    . . .

```

```

23     train_loss /= len(train_iter)
24     val_loss = 0.0
25     model.eval()
26     pbar = tqdm(enumerate(val_iter), total=len(val_iter), leave=False)
27     pbar.set_description(f"Epoch {epoch}")
28     for it, batch in pbar:
29         with torch.no_grad():
30             outputs = model(batch.text[0], batch.text[1].cpu())
31             loss = loss_func(outputs, batch.label.unsqueeze(1)).cpu()
32
33             val_loss += loss
34
35     val_loss /= len(val_iter)
36     if val_loss < min_loss:
37         min_loss = val_loss
38         best_model = model.state_dict()
39     else:
40         cur_patience += 1
41         if cur_patience == patience:
42             cur_patience = 0
43             break
44
45     print('Epoch: {}, Training Loss: {}, Validation Loss: {}'.format(epoch, train_loss, val_loss))
46     model.load_state_dict(best_model)
47
48 def freeze_embeddings(model, req_grad=False):
49     for c_p in model.embedding.parameters():
50         c_p.requires_grad = req_grad
51
52 def train_cnn(model, loss_func, train_iter, val_iter, max_epochs, patience, max_grad_norm=2, num_freeze_epochs=0):
53     min_loss = np.inf
54
55     cur_patience = 0
56
57     freeze_embeddings(model)
58
59     for epoch in range(1, max_epochs + 1):
60         train_loss = 0.0
61         model.train()
62         pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
63         pbar.set_description(f"Epoch {epoch}")
64
65         if epoch > num_freeze_epochs:
66             freeze_embeddings(model, True)
67
68         for it, batch in pbar:
69             opt.zero_grad()
70
71             outputs = model(batch.text)
72             loss = loss_func(outputs, batch.label.unsqueeze(1))
73             loss.backward()
74             if max_grad_norm is not None:
75                 torch.nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
76             opt.step()
77
78             train_loss += loss.cpu().detach()
79
80     train_loss /= len(train_iter)
81     val_loss = 0.0
82     model.eval()
83     pbar = tqdm(enumerate(val_iter), total=len(val_iter), leave=False)
84     pbar.set_description(f"Epoch {epoch}")
85     for it, batch in pbar:
86         with torch.no_grad():
87             outputs = model(batch.text)
88             loss = loss_func(outputs, batch.label.unsqueeze(1)).cpu()
89
90             val_loss += loss
91
92     val_loss /= len(val_iter)
93     if val_loss < min_loss:
94         min_loss = val_loss
95         best_model = model.state_dict()
96     else:
97         cur_patience += 1
98         if cur_patience == patience:
99             cur_patience = 0
100            break
101
102     print('Epoch: {}, Training Loss: {}, Validation Loss: {}'.format(epoch, train_loss, val_loss))
103     model.load_state_dict(best_model)

```

Датасет на котором мы будем проводить эксперименты это комментарии к фильмам из сайта IMDB.

```

1 train, test = datasets.IMDB.splits(TEXT, LABEL) # загрузим датасет
2 train, valid = train.split(random_state=random.seed(SEED)) # разобьем на части

1 TEXT.build_vocab(train)
2 LABEL.build_vocab(train)

1 device = "cuda" if torch.cuda.is_available() else "cpu"
2
3 train_iter, valid_iter, test_iter = BucketIterator.splits(
4     (train, valid, test),
5     batch_size = 64,
6     sort_within_batch = True,
7     device = device)

```

▼ RNN

Для начала попробуем использовать рекуррентные нейронные сети. На семинаре вы познакомились с GRU, вы можете также попробовать LSTM. Можно использовать для классификации как hidden_state, так и output последнего токена.

```

1 class RNNBaseline(nn.Module):
2     def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim, n_layers,
3                  bidirectional, dropout, pad_idx):
4
5         super().__init__()
6
7         self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx = pad_idx)
8
9         self.rnn = nn.LSTM(input_size=embedding_dim, hidden_size=hidden_dim, num_layers=n_layers, bidirectional=bidirectional, dropout=dropout)
10
11         self.dropout = nn.Dropout(dropout)
12
13         self.fc = nn.Linear(2 * hidden_dim, output_dim)
14
15
16     def forward(self, text, text_lengths):
17
18         #text = [sent len, batch size]
19
20         embedded = self.embedding(text)
21
22         embedded = embedded.permute(1, 0, 2)
23
24         hidden = None
25
26         for i in range(text_lengths.max()):
27             hidden = self.rnn(embedded[i], hidden)
28
29         hidden = hidden[-1]
30
31         hidden = hidden.reshape(-1, hidden.size(2))
32
33         output = self.fc(hidden)
34
35         return output

```

```

 2 #embedded = [sent len, batch size, emb dim]
 3
 4 #pack sequence
 5 packed_embedded = nn.utils.rnn.pack_padded_sequence(embedded, text_lengths)
 6
 7 # cell arg for LSTM, remove for GRU
 8 packed_output, (hidden, cell) = self.rnn(packed_embedded)
 9 #unpack sequence
10 output, output_lengths = nn.utils.rnn.pad_packed_sequence(packed_output)
11
12 #output = [sent len, batch size, hid dim * num directions]
13 #output over padding tokens are zero tensors
14
15 #hidden = [num layers * num directions, batch size, hid dim]
16 #cell = [num layers * num directions, batch size, hid dim]
17
18 #concat the final forward (hidden[-2,:,:]) and backward (hidden[-1,:,:]) hidden layers
19 #and apply dropout
20
21 hidden = torch.cat((hidden[-2,:,:], hidden[-1,:,:]), 1)
22 hidden = self.dropout(hidden)
23
24 #hidden = [batch size, hid dim * num directions] or [batch_size, hid dim * num directions]
25
26 return self.fc(hidden)

```

Поиграйтесь с гиперпараметрами

```

1 vocab_size = len(TEXT.vocab)
2 emb_dim = 300
3 hidden_dim = 256
4 output_dim = 1
5 n_layers = 2
6 bidirectional = True
7 dropout = 0.5
8 PAD_IDX = TEXT.vocab.stoi[TEXT.pad_token]
9
10 rnn_model = RNNBaseline(
11     vocab_size=vocab_size,
12     embedding_dim=emb_dim,
13     hidden_dim=hidden_dim,
14     output_dim=output_dim,
15     n_layers=n_layers,
16     bidirectional=bidirectional,
17     dropout=dropout,
18     pad_idx=PAD_IDX
19 ).to(device)

```

```
1 opt = torch.optim.Adam(rnn_model.parameters(), lr=1e-4)
2 loss_func = nn.BCEWithLogitsLoss()
```

Обучите сетку! Используйте любые вам удобные инструменты, Catalyst, PyTorch Lightning или свои велосипеды.

```

1 max_epochs = 50
2 patience=15
3
4 train_rnn(rnn_model, loss_func, train_iter, valid_iter, max_epochs, patience)

Epoch: 1, Training Loss: 0.6742545366287231, Validation Loss: 0.5824978351593018
Epoch: 2, Training Loss: 0.5109234732479858, Validation Loss: 0.4636767051696777
Epoch: 3, Training Loss: 0.4130497827824937, Validation Loss: 0.4198527932167053
Epoch: 4, Training Loss: 0.3488076254219498, Validation Loss: 0.42421379672159574
Epoch: 5, Training Loss: 0.285278138472377, Validation Loss: 0.48121873432942
Epoch: 6, Training Loss: 0.2449284434911753, Validation Loss: 0.425086581574842
Epoch: 7, Training Loss: 0.19758240878528, Validation Loss: 0.43543775147352
Epoch: 8, Training Loss: 0.15538778628101, Validation Loss: 0.4638367526363916
Epoch: 9, Training Loss: 0.125805391000246, Validation Loss: 0.47576326452624
Epoch: 10, Training Loss: 0.09957393807334473, Validation Loss: 0.639642315752724
Epoch: 11, Training Loss: 0.0761903203432242, Validation Loss: 0.55092426665262246
Epoch: 12, Training Loss: 0.06665814753824865, Validation Loss: 0.6938986632152153
Epoch: 13, Training Loss: 0.04886015877126744, Validation Loss: 0.66979893280077
Epoch: 14, Training Loss: 0.041927007352112734, Validation Loss: 0.718866600444301
Epoch: 15, Training Loss: 0.032427367379875, Validation Loss: 0.75761894113769
Epoch: 16, Training Loss: 0.031396606443045044, Validation Loss: 0.8372392479424963
Epoch: 17, Training Loss: 0.022849265449632645, Validation Loss: 0.8494941584762972
Epoch: 18, Training Loss: 0.01542875749376486, Validation Loss: 0.8026856983902167

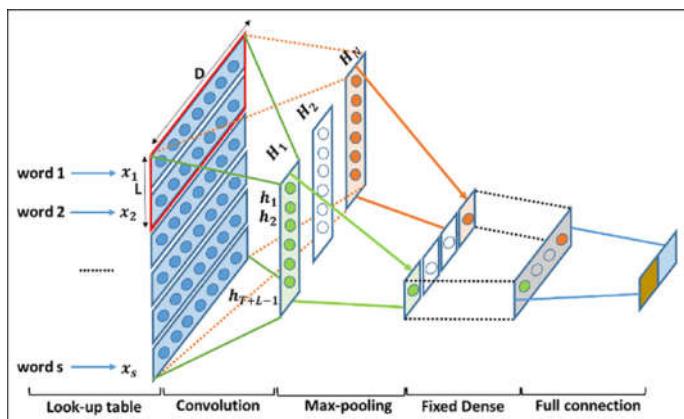
```

Посчитайте f1-score вашего классификатора на тестовом датасете.

Ответ:

```
1 calc_f1_score(rnn_model, test_iter, lambda model, batch: model(batch.text[0], batch.text[1].cpu()))
```

▼ CNN



Для классификации текстов также часто используют сверточные нейронные сети. Идея в том, что как правило сентимент содержит словосочетания из двух-трех слов, например "очень хороший фильм" или "невероятная скуча". Проходясь сверткой по этим словам мы получим какой-то большой скор и выхватим его с помощью MaxPool. Далее идет обычная полно связная сетька. Важный момент: свортки применяются не последовательно, а параллельно. Давайте попробуем!

```
1 TEXT = Field(sequential=True, lower=True, batch_first=True) # batch_first тк мы используем conv  
2 LABEL = LabelField(batch_first=True, dtype=torch.float)
```

```

3
4 train, tst = datasets.IMDB.splits(TEXT, LABEL)
5 trn, vld = train.split(random_state=random.seed(SEED))
6
7 TEXT.build_vocab(trn)
8 LABEL.build_vocab(trn)
9
10 device = "cuda" if torch.cuda.is_available() else "cpu"
11
12 train_iter, val_iter, test_iter = BucketIterator.splits(
13     (trn, vld, tst),
14     batch_sizes=(128, 256, 256),
15     sort=False,
16     sort_key= lambda x: len(x.src),
17     sort_within_batch=False,
18     device=device,
19     repeat=False,
20 )

```

Вы можете использовать Conv2d с in_channels=1, kernel_size=(kernel_sizes[0], emb_dim)) или Conv1d с in_channels=emb_dim, kernel_size=kernel_size[0]. Но хорошенько подумайте над shape в обоих случаях.

```

1 class CNN(nn.Module):
2     def __init__(self, vocab_size, emb_dim, out_channels, kernel_sizes, dropout=0.5,
3      ):
4         super().__init__()
5
6         self.embedding = nn.Embedding(vocab_size, emb_dim)
7         self.conv_0 = nn.Conv1d(in_channels=emb_dim, out_channels=out_channels, kernel_size=kernel_sizes[0])
8
9         self.conv_1 = nn.Conv1d(in_channels=emb_dim, out_channels=out_channels, kernel_size=kernel_sizes[1])
10
11        self.conv_2 = nn.Conv1d(in_channels=emb_dim, out_channels=out_channels, kernel_size=kernel_sizes[2])
12
13        self.fc = nn.Linear(len(kernel_sizes) * out_channels, 1)
14
15        self.dropout = nn.Dropout(dropout)
16
17
18    def forward(self, text):
19        embedded = self.embedding(text)
20
21        embedded = embedded.permute(0, 2, 1)
22
23        conved_0 = F.relu(self.conv_0(embedded)) # may be reshape here
24        conved_1 = F.relu(self.conv_1(embedded)) # may be reshape here
25        conved_2 = F.relu(self.conv_2(embedded)) # may be reshape here
26
27        pooled_0 = F.max_pool1d(conved_0, conved_0.shape[2]).squeeze(2)
28        pooled_1 = F.max_pool1d(conved_1, conved_1.shape[2]).squeeze(2)
29        pooled_2 = F.max_pool1d(conved_2, conved_2.shape[2]).squeeze(2)
30
31        cat = self.dropout(torch.cat((pooled_0, pooled_1, pooled_2), dim=1))
32
33        return self.fc(cat)
34
35
36
37
38
39

```

```

1 kernel_sizes = [3, 4, 5]
2 vocab_size = len(TEXT.vocab)
3 out_channels=64
4 dropout = 0.5
5 dim = 300
6
7 cnn_model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=out_channels, kernel_sizes=kernel_sizes, dropout=dropout).to(device)

```

```

1 opt = torch.optim.Adam(cnn_model.parameters(), lr=1e-4)
2 loss_func = nn.BCEWithLogitsLoss()

```

Обучите!

```

1 max_epochs = 50
2 patience = 15
3
4 train_cnn(cnn_model, loss_func, train_iter, val_iter, max_epochs, patience)

```

```
Epoch: 1, Training Loss: 0.7576324343681335, Validation Loss: 0.6517045497894278
Epoch: 2, Training Loss: 0.6899340748786926, Validation Loss: 0.6062830686569214
Epoch: 3, Training Loss: 0.6387899518013, Validation Loss: 0.5613199472427368
Epoch: 4, Training Loss: 0.5910921692848206, Validation Loss: 0.5337907075881958
```

Посчитайте f1-score вашего классификатора.

Ответ:

```
Epoch: 9, Training Loss: 0.4728925824165344, Validation Loss: 0.4650088548660278
1 calc_f1_score(cnn_model, test_iter, lambda model, batch: model(batch.text))
0.8496467471122742
Epoch: 14, Training Loss: 0.4180844068//51//, Validation Loss: 0.4302950501441956
```

▼ Интерпретируемость

Посмотрим, куда смотрит наша модель. Достаточно запустить код ниже.

```
Epoch: 20, Training Loss: 0.39227400000000004, Validation Loss: 0.3502274477527474
1 !pip install -q captum
Epoch: 20, Training Loss: 0.39227400000000004, Validation Loss: 0.3502274477527474
from captum.attr import LayerIntegratedGradients, TokenReferenceBase, visualization
1
2 PAD_IND = TEXT.vocab.stoi['pad']
3
4 token_reference = TokenReferenceBase(reference_token_idx=PAD_IND)
5 lig = LayerIntegratedGradients(cnn_model, cnn_model.embedding)
6
7 Epoch: 32. Training Loss: 0.26171180605888367. Validation Loss: 0.35848888754844666
8
9 # accumulate couple samples in this array for visualization purposes
10 vis_data_records_ig = []
11
12 def interpret_sentence(model, sentence, min_len = 7, label = 0):
13     model.eval()
14     text = [tok for tok in TEXT.tokenize(sentence)]
15     if len(text) < min_len:
16         text += ['pad'] * (min_len - len(text))
17     indexed = [TEXT.vocab.stoi[t] for t in text]
18
19     model.zero_grad()
20
21     input_indices = torch.tensor(indexed, device=device)
22     input_indices = input_indices.unsqueeze(0)
23
24     # input_indices dim: [sequence_length]
25     seq_length = min_len
26
27     # predict
28     pred = forward_with_sigmoid(model, input_indices).item()
29     pred_ind = round(pred)
30
31     # generate reference indices for each sample
32     reference_indices = token_reference.generate_reference(seq_length, device=device).unsqueeze(0)
33
34     # compute attributions and approximation delta using layer integrated gradients
35     attributions_ig, delta = lig.attribute(input_indices, reference_indices, \
36                                             n_steps=5000, return_convergence_delta=True)
37
38     print('pred: ', LABEL.vocab.itos[pred_ind], '(', '%.2f'%pred, ')', ', delta: ', abs(delta))
39
40     add_attributions_to_visualizer(attributions_ig, text, pred, pred_ind, label, delta, vis_data_records_ig)
41
42 def add_attributions_to_visualizer(attributions, text, pred, pred_ind, label, delta, vis_data_records):
43     attributions = attributions.sum(dim=2).squeeze(0)
44     attributions = attributions / torch.norm(attributions)
45     attributions = attributions.cpu().detach().numpy()
46
47     # storing couple samples in an array for visualization purposes
48     vis_data_records.append(visualization.VisualizationDataRecord(
49         attributions,
50         pred,
51         LABEL.vocab.itos[pred_ind],
52         LABEL.vocab.itos[label],
53         LABEL.vocab.itos[1],
54         attributions.sum(),
55         text,
56         delta))
57
58
1 interpret_sentence(cnn_model, 'It was a fantastic performance !', label=1)
2 interpret_sentence(cnn_model, 'Best film ever', label=1)
3 interpret_sentence(cnn_model, 'Such a great show!', label=1)
4 interpret_sentence(cnn_model, 'It was a horrible movie', label=0)
5 interpret_sentence(cnn_model, 'I\'ve never watched something as bad', label=0)
6 interpret_sentence(cnn_model, 'It is a disgusting movie!', label=0)

pred: pos ( 0.99 ), delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.30 ), delta: tensor([7.6005e-06], device='cuda:0', dtype=torch.float64)
pred: pos ( 0.96 ), delta: tensor([1.3040e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.89 ), delta: tensor([3.4205e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.12 ), delta: tensor([7.6346e-05], device='cuda:0', dtype=torch.float64)
pred: pos ( 0.69 ), delta: tensor([1.7903e-05], device='cuda:0', dtype=torch.float64)
```

Попробуйте добавить свои примеры!

```
1 print('Visualize attributions based on Integrated Gradients')
2 visualization.visualize_text(vis_data_records_ig)
```

Visualize attributions based on Integrated Gradients

Legend: Negative Neutral Positive

True Label Predicted Label Attribution Label Attribution Score Word Importance

▼ ЭМБЭДИНГИ СЛОВ

Вы ведь не забыли, как мы можем применить знания о word2vec и GloVe. Давайте попробуем!

```

1 nea     nea (0.09)    pos      -0.45      It was a horrible movie nad nad
1 TEXT.build_vocab(trn, vectors='glove.840B.300d')
2 LABEL.build_vocab(trn)
3
4 word_embeddings = TEXT.vocab.vectors
5
6 train, tst = datasets.IMDB.splits(TEXT, LABEL)
7 trn, vld = train.split(random_state=random.seed(SEED))
8
9 device = "cuda" if torch.cuda.is_available() else "cpu"
10
11 train_iter, val_iter, test_iter = BucketIterator.splits(
12     (trn, vld, tst),
13     batch_sizes=(128, 256, 256),
14     sort=False,
15     sort_key= lambda x: len(x.src),
16     sort_within_batch=False,
17     device=device,
18     repeat=False,
19 )
20
21 kernel_sizes = [3, 4, 5]
22 vocab_size = len(TEXT.vocab)
23 out_channels=64
24 dropout = 0.5
25 dim = 300
26
27 cnn_and_embeddings_model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=out_channels, kernel_sizes=kernel_sizes, dropout=dropout)
28
29 word_embeddings = TEXT.vocab.vectors
30
31 prev_shape = cnn_and_embeddings_model.embedding.weight.shape
32
33 cnn_and_embeddings_model.embedding.weight.data = torch.clone(word_embeddings)
34
35 assert prev_shape == cnn_and_embeddings_model.embedding.weight.shape
36
37 cnn_and_embeddings_model.to(device)
38
39 CNN(
40     (embedding): Embedding(202268, 300)
41     (conv_0): Conv1d(300, 64, kernel_size=(3,), stride=(1,))
42     (conv_1): Conv1d(300, 64, kernel_size=(4,), stride=(1,))
43     (conv_2): Conv1d(300, 64, kernel_size=(5,), stride=(1,))
44     (fc): Linear(in_features=192, out_features=1, bias=True)
45     (dropout): Dropout(p=0.5, inplace=False)
46 )
47
48 opt = torch.optim.Adam(cnn_and_embeddings_model.parameters(), lr=1e-4)
49 loss_func = nn.BCEWithLogitsLoss()

```

Вы знаете, что делать.

```

1 max_epochs = 50
2 patience = 15
3
4 train_cnn(cnn_and_embeddings_model, loss_func, train_iter, val_iter, max_epochs, patience, num_freeze_epochs=0)

Epoch: 1, Training Loss: 0.6762555241584778, Validation Loss: 0.6392387877540588
Epoch: 2, Training Loss: 0.6033423542976379, Validation Loss: 0.5459826733833313
Epoch: 3, Training Loss: 0.49212383376721, Validation Loss: 0.45212383376721
Epoch: 4, Training Loss: 0.4266327032646646, Validation Loss: 0.3895754474482896
Epoch: 5, Training Loss: 0.38875171542167664, Validation Loss: 0.3815528611190796
Epoch: 6, Training Loss: 0.3620424508465393, Validation Loss: 0.365334379196167
Epoch: 7, Training Loss: 0.33915168046951204, Validation Loss: 0.3537197789083557
Epoch: 8, Training Loss: 0.31842857509258423, Validation Loss: 0.3433770537736494
Epoch: 9, Training Loss: 0.30093020198047485, Validation Loss: 0.33546820382936906
Epoch: 10, Training Loss: 0.289275732265876, Validation Loss: 0.3285302517978279
Epoch: 11, Training Loss: 0.2626259326938145, Validation Loss: 0.3221765766070557
Epoch: 12, Training Loss: 0.2462836372761917, Validation Loss: 0.3174632489681244
Epoch: 13, Training Loss: 0.2398123499748074, Validation Loss: 0.31211230158805847
Epoch: 14, Training Loss: 0.2127387821674347, Validation Loss: 0.3088242789636688
Epoch: 15, Training Loss: 0.19626827538013458, Validation Loss: 0.30575281381607856
Epoch: 16, Training Loss: 0.18307611346244812, Validation Loss: 0.3018568999385834
Epoch: 17, Training Loss: 0.16972137987613678, Validation Loss: 0.2993110418319702
Epoch: 18, Training Loss: 0.15460239350795746, Validation Loss: 0.29740708125694275
Epoch: 19, Training Loss: 0.14086130261421204, Validation Loss: 0.2963598966598511
Epoch: 20, Training Loss: 0.1289202570915222, Validation Loss: 0.29563918709754944
Epoch: 21, Training Loss: 0.11790939420461655, Validation Loss: 0.2952145840035248
Epoch: 22, Training Loss: 0.105301462111385727, Validation Loss: 0.2947268486022949
Epoch: 23, Training Loss: 0.0947084494089127, Validation Loss: 0.2952643036842346
Epoch: 24, Training Loss: 0.085772204495668411, Validation Loss: 0.29551824972733017
Epoch: 25, Training Loss: 0.075457978571518, Validation Loss: 0.2977536206863403
Epoch: 26, Training Loss: 0.06945324689149857, Validation Loss: 0.2990273232283936
Epoch: 27, Training Loss: 0.06282911585569382, Validation Loss: 0.3005805015653965
Epoch: 28, Training Loss: 0.055815866629081726, Validation Loss: 0.30180710554122925
Epoch: 29, Training Loss: 0.048888778458782, Validation Loss: 0.30486446619633813
Epoch: 30, Training Loss: 0.043837517499923766, Validation Loss: 0.3083876967438115
Epoch: 31, Training Loss: 0.03895726986669815, Validation Loss: 0.3083434184919136
Epoch: 32, Training Loss: 0.03586233915608074, Validation Loss: 0.3109781282855144
Epoch: 33, Training Loss: 0.032937191438652927, Validation Loss: 0.31365652238845923
Epoch: 34, Training Loss: 0.02762993258863292, Validation Loss: 0.31838181614875793
Epoch: 35, Training Loss: 0.025365378767647324, Validation Loss: 0.3211582899903628
Epoch: 36, Training Loss: 0.022864093506375748, Validation Loss: 0.32449501752853394

```

Посчитайте f1-score вашего классификатора.

Ответ:

```

1 calc_f1_score(cnn_and_embeddings_model, test_iter, lambda model, batch: model(batch.text))
2
3 0.872462272644043

```

Проверим насколько все хорошо!

```

1 PAD_IND = TEXT.vocab.stoi['pad']
2
3 token_reference = TokenReferenceBase(reference_token_idx=PAD_IND)
4 lIG = LayerIntegratedGradients(cnn_and_embeddings_model, cnn_and_embeddings_model.embedding)
5 vis_data_records_lIG = []
6
7 interpret_sentence(cnn_and_embeddings_model, 'It was a fantastic performance !', label=1)
8 interpret_sentence(cnn_and_embeddings_model, 'Best film ever', label=1)
9 interpret_sentence(cnn_and_embeddings_model, 'Such a great show!', label=1)
10 interpret_sentence(cnn_and_embeddings_model, 'It was a horrible movie', label=0)
11 interpret_sentence(cnn_and_embeddings_model, 'I've never watched something as bad', label=0)
12 interpret_sentence(cnn_and_embeddings_model, 'It is a disgusting movie!', label=0)

```

```

pred: pos ( 0.98 ), delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.01 ), delta: tensor([2.585e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.35 ), delta: tensor([0.0002], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ), delta: tensor([5.477e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.43 ), delta: tensor([0.0002], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ), delta: tensor([1.0463e-05], device='cuda:0', dtype=torch.float64)

```

```
1 print('Visualize attributions based on Integrated Gradients')
2 visualization.visualize_text(vis_data_records_ig)
```

Visualize attributions based on Integrated Gradients

Legend: Negative Neutral Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (0.98)	pos	1.85	It was a fantastic performance I pad
pos	neg (0.01)	pos	1.12	Best film ever pad pad pad
pos	neg (0.35)	pos	1.57	Such a great show! pad pad pad
neg	neg (0.00)	pos	0.06	It was a horrible movie pad pad
neg	neg (0.43)	pos	1.65	I've never watched something as bad pad
neg	neg (0.00)	pos	0.12	It is a disgusting movie! pad pad

Legend: Negative Neutral Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (0.98)	pos	1.85	It was a fantastic performance I pad
pos	neg (0.01)	pos	1.12	Best film ever pad pad pad
pos	neg (0.35)	pos	1.57	Such a great show! pad pad pad
neg	neg (0.00)	pos	0.06	It was a horrible movie pad pad
neg	neg (0.43)	pos	1.65	I've never watched something as bad pad
neg	neg (0.00)	pos	0.12	It is a disgusting movie! pad pad

▼ CNN + Embeddings + Frozen N first epochs

```

1 cnn_and_frozen_embeddings_model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=out_channels, kernel_sizes=kernel_sizes, dropout
2
3 word_embeddings = TEXT.vocab.vectors
4
5 prev_shape = cnn_and_frozen_embeddings_model.embedding.weight.shape
6
7 cnn_and_frozen_embeddings_model.embedding.weight.data = torch.clone(word_embeddings)
8
9 assert prev_shape == cnn_and_frozen_embeddings_model.embedding.weight.shape
10 cnn_and_frozen_embeddings_model.to(device)
11
12 opt = torch.optim.Adam(cnn_and_frozen_embeddings_model.parameters(), lr=1e-4)
13 loss_func = nn.BCEWithLogitsLoss()
14
15 max_epochs = 50
16 patience = 15
17
18 train_cnn(cnn_and_frozen_embeddings_model, loss_func, train_iter, val_iter, max_epochs, patience, num_freeze_epochs=20)

Epoch: 1, Training Loss: 0.676727958729675, Validation Loss: 0.6421320824431458
Epoch: 2, Training Loss: 0.614164385795593, Validation Loss: 0.565899919985469
Epoch: 3, Training Loss: 0.526130672605312, Validation Loss: 0.4755525731598306
Epoch: 4, Training Loss: 0.458656686630249, Validation Loss: 0.42740555140398513
Epoch: 5, Training Loss: 0.42121160705200805, Validation Loss: 0.4015980660910375
Epoch: 6, Training Loss: 0.39548927515411377, Validation Loss: 0.38501882553109586
Epoch: 7, Training Loss: 0.37871563434680883, Validation Loss: 0.37361207604498264
Epoch: 8, Training Loss: 0.3680460158348883, Validation Loss: 0.365896444985198975
Epoch: 9, Training Loss: 0.35644228981407166, Validation Loss: 0.357986661338886
Epoch: 10, Training Loss: 0.3496182858943930, Validation Loss: 0.3519825941294998
Epoch: 11, Training Loss: 0.3371342428717346, Validation Loss: 0.3468288481235984
Epoch: 12, Training Loss: 0.3277920186519623, Validation Loss: 0.34187641739845276
Epoch: 13, Training Loss: 0.320948898794226685, Validation Loss: 0.3384256958961487
Epoch: 14, Training Loss: 0.3118315041065216, Validation Loss: 0.3342123793601099
Epoch: 15, Training Loss: 0.30361726880035, Validation Loss: 0.3307430148124695
Epoch: 16, Training Loss: 0.29895299673080444, Validation Loss: 0.32824888545184326
Epoch: 17, Training Loss: 0.292708158493042, Validation Loss: 0.325955945223236
Epoch: 18, Training Loss: 0.28394633531570435, Validation Loss: 0.3225109875202179
Epoch: 19, Training Loss: 0.28482427833801126, Validation Loss: 0.328369567680359
Epoch: 20, Training Loss: 0.2723568379879996, Validation Loss: 0.31801366806030273
Epoch: 21, Training Loss: 0.26439484895388031, Validation Loss: 0.31352129578590393
Epoch: 22, Training Loss: 0.2518346309661865, Validation Loss: 0.3099350167274475
Epoch: 23, Training Loss: 0.23830914497375488, Validation Loss: 0.307134687905432
Epoch: 24, Training Loss: 0.223856122606363, Validation Loss: 0.3037700857029724
Epoch: 25, Training Loss: 0.2122745968958154, Validation Loss: 0.30805315661430359
Epoch: 26, Training Loss: 0.1973966957659094, Validation Loss: 0.297966063322029114
Epoch: 27, Training Loss: 0.18485459685325623, Validation Loss: 0.2954931887689667
Epoch: 28, Training Loss: 0.17324137464057415, Validation Loss: 0.29324138164520264
Epoch: 29, Training Loss: 0.161745124657833, Validation Loss: 0.2841251112779493
Epoch: 30, Training Loss: 0.1505454078601282, Validation Loss: 0.2801380022386577
Epoch: 31, Training Loss: 0.1379713409895406, Validation Loss: 0.28846704959569385
Epoch: 32, Training Loss: 0.12843388965434986, Validation Loss: 0.2872713041138554
Epoch: 33, Training Loss: 0.11015233731269836, Validation Loss: 0.2870180016795593
Epoch: 34, Training Loss: 0.10873080798042877, Validation Loss: 0.28550891351699829
Epoch: 35, Training Loss: 0.10037156194448471, Validation Loss: 0.2840944325685911
Epoch: 36, Training Loss: 0.09172337532351125, Validation Loss: 0.2848204411506653
Epoch: 37, Training Loss: 0.08283637464046478, Validation Loss: 0.28615668416023254
Epoch: 38, Training Loss: 0.07585844397544861, Validation Loss: 0.284568373136139
Epoch: 39, Training Loss: 0.07033748179674149, Validation Loss: 0.287023663529813
Epoch: 40, Training Loss: 0.06372224539518356, Validation Loss: 0.285308508682251
Epoch: 41, Training Loss: 0.05621589999936789, Validation Loss: 0.288138747215271
Epoch: 42, Training Loss: 0.052678531919237518, Validation Loss: 0.2900582551956177
Epoch: 43, Training Loss: 0.04726078733801842, Validation Loss: 0.2891470193862915
Epoch: 44, Training Loss: 0.0423966683447361, Validation Loss: 0.29121874713169
Epoch: 45, Training Loss: 0.037734924665607986, Validation Loss: 0.29307987819747925
Epoch: 46, Training Loss: 0.034464605152606964, Validation Loss: 0.2945287525653839
Epoch: 47, Training Loss: 0.031634438782930374, Validation Loss: 0.2983073592185974
Epoch: 48, Training Loss: 0.028561849147081375, Validation Loss: 0.298458993652344
Epoch: 49, Training Loss: 0.02543807215988636, Validation Loss: 0.30106598138809204
Epoch: 50, Training Loss: 0.02351159229874611, Validation Loss: 0.3031412959098816

```

```
1 calc_f1_score(cnn_and_frozen_embeddings_model, test_iter, lambda model, batch: model(batch.text))

0.8806599974632263
```

Score для всех экспериментов:

- RNN (LSTM) - 0.8374814987182617
- CNN - 0.8496467471122742
- CNN + Glove Embeddings - 0.872462272644043
- CNN + Glove Embeddings (frozen first N epochs) - 0.8806599974632263

Выходы

- CNN работает лучше чем RNN
- Предобученные эмбеддинги помогают улучшить результат
- Замораживание весов эмбеддингов на первые N эпохи помогает улучшить результат еще лучше

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.